SMART LIBRARIAN - COMPLETE TEAM DEMO TEST DECK

Full four-dataset live-LLM test run

This is a manual test plan, not a record of tests executed today. Run every dataset chapter, then the shared checks. Record actual results separately. The four datasets are Commercial, Supply, Market Access, and Launch from the official synthetic challenge builder pack, not real Bayer performance. Each numbered section is a slide in the team testing deck.

44 test sections | 4 datasets x 5 steps + shared checks | Manual test plan - not today's results

Arrow keys / Space navigate. The Contents button jumps to any section. Print produces one landscape page per section.

Orange = expected result Teal = evidence to record Dot = action / instruction

SETUP & GROUND RULES

01 Mission: test the whole journey

- Test all four datasets with live LLM extraction, explicit form confirmation, ownership, approval, frozen reports, downloads, saved plans, and REUSE.
- Also exercise steward refusal, rejection, scope refusal, unsupported intake, and honest offline fallback.
- The UI, API, and exported reports run locally. Optional LLM inference is remote.
- These are simulated actors and approvals, not real authentication. Never expose the API remotely.

SETUP & GROUND RULES

02 Run order and responsibilities

- One tester drives the browser; a second records evidence. The approver in the UI is still a simulated actor.
- Run setup first, Commercial C1-C5, Supply S1-S5, Market Access M1-M5, Launch L1-L5, then shared checks G1-G14.
- Use one fresh test database per tester or coordinated run. Do not let concurrent testers approve the same scope unless testing the race scenario.
- Keep a record per case: date, tester, dataset, model/provider, extraction mode, original plan ID, reused plan ID, revision/spec hash, actual result, evidence path, and notes.

SETUP & GROUND RULES

03 First-time setup in PowerShell

- Open the repository folder in File Explorer; type powershell in its address bar and press Enter. Python 3.11 or later is required.

```
python -m venv .venv
.\.venv\Scripts\python.exe -m pip install -e ".[ui,llm,dev]"
```

- Installation needs internet. Calling the venv Python directly avoids activation and execution-policy changes.
- The workbook and frozen evidence ship with the repo. Do not edit or regenerate them. Do not overwrite the existing .env.

SETUP & GROUND RULES

04 Choose fresh approval or existing-state REUSE

- Existing session: try UI http://127.0.0.1:8501 and API http://127.0.0.1:8000. Reuse it for existing-state testing if it responds.
- For every fresh approval step, launch a separate test run with a new database and free ports. Run these commands in the same PowerShell:

```
$env:RUNTIME_DB_PATH = "var/team-demo-$([guid]::NewGuid().ToString('N')).sqlite"
.\.venv\Scripts\python.exe run_demo.py --mode challenge --api-port 8021 --ui-port 8521
```

- This creates a new DB, preserving var/runtime.sqlite. Record the chosen DB path privately before stopping. Never put the DB under data/.
- If either port is busy, choose both new ports, e.g. 8022/8522. Never kill unknown processes. Do not use --background for this isolated test recipe.

SETUP & GROUND RULES

05 Open the right application

- Wait for API ready and UI ready in the launching terminal; keep that terminal open.
- Fresh-run UI: http://127.0.0.1:8521. API health: http://127.0.0.1:8021/api/health. Interactive API docs: http://127.0.0.1:8021/docs.
- Sidebar Simulated requester identity: DEMO-CHALLENGE-REQUESTER. Workspace: Builder pack reports.
- Confirm the synthetic challenge banner and the four Source domain choices. A public/legacy app is the wrong mode for these tests.

SETUP & GROUND RULES

06 Configure live LLM without sharing secrets

- Use the existing approved provider configuration; ask your project owner for access if missing. Never paste a token into slides, chat, request text, screenshots, or test results.
- Aitta configuration names: LLM_PROVIDER=aitta, LLM_MODEL=openai/gpt-oss-120b, AITTA_API_TOKEN. The token value stays private.
- An OpenAI configuration instead uses LLM_PROVIDER=openai and OPENAI_API_KEY with an appropriate configured model. Do not mix credentials.
- Process environment overrides .env; launch the app after configuration. GET /api/capabilities exposes provider configuration status without the token, but does not prove inference works.
- Internet and a valid credential are required for live extraction. Request text and catalog metadata are sent remotely, not source business figures. Never enter confidential text.

SETUP & GROUND RULES

07 What counts as live extraction evidence

- For each dataset, tick Send request text + catalog metadata to the LLM, enter the exact request below, then click Extract intent.

> **EVIDENCE**
> Capture the UI caption Selected mode: live LLM | last extraction: llm. A checked box alone is not evidence.

- Compare the extracted form to the expected scope before editing. Record missing/wrong fields as extraction defects; record manual corrections separately.
- Do not press Prefill during the live extraction test: it replaces extraction state and is not proof that the LLM understood the request.
- offline_fallback_after_llm_unavailable is an honest fallback, NOT a live-LLM pass. Record live testing as blocked and test fallback separately.
- Extraction itself must not create or approve a plan. Existing plan state from an earlier action is not a new extraction result.

COMMERCIAL C1-C5

08 C1: Commercial - exact LLM request

- Workspace Builder pack reports; Source domain Commercial; requester DEMO-CHALLENGE-REQUESTER.

```
Show monthly Nubeqa net sales and sales forecast in Germany for January through March 2026.
```

- Paste into Describe the report, tick the LLM opt-in, click Extract intent, and wait for completion.

EXPECTED

Expected: last extraction: llm. Save a screenshot of the request, caption, and any warnings before changing the form.

COMMERCIAL C1-C5

09 C2: Commercial - confirm exact scope

- Products: Nubeqa (P-NUB). Markets: Germany (DE).
- Source fields: net_sales_m_eur_syn and forecast_m_eur_syn.
- Months: 2026-01-01, 2026-02-01, 2026-03-01. Grain: month.
- Inspect Guided confirmation (authoritative). Record any mismatch, then explicitly correct it if continuing.
- Accountable owner: choose active OWN-001 yourself; no default is selected. Rationale: Team Commercial live LLM test.

EXPECTED

Expected: Create plan (requires approval) becomes enabled. Capture the confirmed form.

COMMERCIAL C1-C5

10 C3: Commercial - refuse steward, then approve

- Click Create plan (requires approval). Record plan ID, revision, spec hash, and status. Fresh DB expected: pending.
- If REUSE appears, use C5 and record fresh approval/steward checks as not exercised on this plan; use a fresh test DB for those checks.
- In Approval (separate actor), choose DEMO-CHALLENGE-STEWARD and click Approve.

> **EXPECTED**
> Expected: API error 403; steward is not an approver; plan remains pending. Capture refusal.

- In the SAME block choose DEMO-CHALLENGE-APPROVER and click Approve. Expected: approved. Capture status and Audit trail.
- Do not switch the sidebar requester identity. It clears current UI plan state.

11 C4: Commercial - render, inspect, download

- Click Render frozen report. Expected: rendered; monthly chart plus frozen source-row table and Audit trail.

> **EXPECTED**
> Expected scope: Nubeqa/Germany, Jan-Mar 2026; three source rows, million EUR units, no aggregation. January source values: net sales 19.66, forecast 18.42; value_type forecast_like_synthetic. These are synthetic, not actual performance.

- Download ALL five: HTML, CSV, ZIP bundle, Power Query kit (CSV+M), Power BI project (PBIP ZIP).

> **EVIDENCE**
> Save under a Commercial-specific evidence folder; export names repeat across domains, so avoid overwriting.

- Open HTML and CSV; compare selected fields, dates, values, units, and source provenance with the frozen preview. Inspect ZIP/kit contents; follow G2-G3 for Power BI.

12 C5: Commercial - prove repeat-request REUSE

- After rendering, record original plan ID and spec hash. Keep the same requester and owner OWN-001.
- Re-enter the exact C1 request with LLM enabled, click Extract intent, and capture last extraction: llm again.
- Confirm EXACTLY the C2 product, market, fields, months, grain, and owner; click Create plan (requires approval).

EXPECTED

Expected: equivalent approved scope already exists / REUSE. Returned plan ID and spec hash match the original; status is rendered and the original frozen result is available.

- No new pending plan or new approval is needed. Re-download CSV and compare its source rows with the first CSV; do not use volatile ZIP bytes as a reuse test.

EVIDENCE

Capture original and reused IDs, the REUSE notice, and Audit trail. Keep source fingerprint unchanged.

13 S1: Supply - exact LLM request

- Return to Builder pack reports; choose Source domain Supply; keep DEMO-CHALLENGE-REQUESTER.

```
Show a current snapshot of Nubeqa in Germany: on-time in-full, inventory in weeks, and supply risk status.
```

- Enter Describe the report, ensure LLM opt-in is checked, click Extract intent.

EXPECTED

Expected: last extraction: llm. Capture request/caption/warnings before edits. Do not use the Prefill button as extraction evidence.

14 S2: Supply - confirm snapshot scope

- Products P-NUB; Markets DE.
- Source fields otif_pct_syn, inventory_weeks_syn, risk_status.
- Grain snapshot; months empty. Expected caption: Snapshot source: no period selection applies.
- There must be no invented dated trend. Record extraction mismatches before correcting the authoritative form.
- Explicit owner OWN-001; rationale Team Supply live LLM test. Capture completed form.

15 S3: Supply - pending, steward refusal, approval

- Click Create plan (requires approval); record ID, revision, hash. Fresh DB expected: pending. If REUSE, follow S5 and mark fresh-only checks not exercised.
- Approval block: choose DEMO-CHALLENGE-STEWARD, click Approve. Expected API error 403 and unchanged pending status.
- Choose DEMO-CHALLENGE-APPROVER in that block; click Approve. Expected approved.
- Keep the sidebar identity unchanged. Capture pending, refusal, approved, and Audit trail.

16 S4: Supply - frozen snapshot and all exports

- Click Render frozen report. Expected rendered; snapshot bars for numeric fields and a table preserving categorical risk_status.
- For P-NUB/DE, source values include otif_pct_syn 87.9 and inventory_weeks_syn 8.8. Preserve the source risk label and native units; no invented monthly series or totals.
- Download HTML, CSV, ZIP bundle, Power Query kit (CSV+M), Power BI project (PBIP ZIP) into a Supply-specific folder.
- Open HTML/CSV, compare to preview, inspect ZIP/kit, and perform G2-G3 for this domain's PBIP.

EVIDENCE

Capture table, chart, Audit trail, and evidence filenames.

17 S5: Supply - repeat live request and REUSE

EVIDENCE
Record rendered Supply plan ID/hash; repeat S1 with LLM opt-in and capture last extraction: llm.

- Confirm P-NUB/DE; otif_pct_syn, inventory_weeks_syn, risk_status; empty months; snapshot; SAME owner OWN-001.
- Click Create plan (requires approval). Expected REUSE of original plan ID/hash, rendered status, original frozen result.
- No fresh approval needed. Re-download CSV and compare source rows. Capture notice, original/reused IDs, and Audit trail.
- A dated period or changed owner is not this identical-scope test.

18 M1: Market Access - exact LLM request

- Builder pack reports; Source domain Market Access; requester DEMO-CHALLENGE-REQUESTER.

```
Show the current market access snapshot for Nubeqa in Germany: coverage percentage and reimbursement status.
```

- Enter request, tick LLM opt-in, click Extract intent; capture last extraction: llm and any warnings before edits.
- The request is for a current snapshot, not a month-by-month trend.

19 M2: Market Access - confirm exact form

- Products P-NUB; Markets DE.
- Source fields coverage_pct_syn and reimbursement_status.
- Grain snapshot; months empty; no period selector applies. Record mismatches before manual correction.
- Explicit owner OWN-001; rationale Team Market Access live LLM test.

EXPECTED
Expected Create plan enabled. Capture completed authoritative form.

20 M3: Market Access - governance checks

- Create plan; record ID/hash/revision. Fresh DB expected pending; existing REUSE means fresh-only checks are not exercised.
- Choose DEMO-CHALLENGE-STEWARD in Approval (separate actor); Approve must show API error 403, status still pending.
- Choose DEMO-CHALLENGE-APPROVER in the same block; Approve must yield approved.
- Keep DEMO-CHALLENGE-REQUESTER in the sidebar. Capture pending/refusal/approval and Audit trail.

21 M4: Market Access - render and export

- Click Render frozen report. Expected rendered snapshot chart/table; coverage_pct_syn is 65.6 for P-NUB/DE.
- Preserve reimbursement_status as a category, not a number. No invented historical series or averaged percentages.
- Download HTML, CSV, ZIP bundle, Power Query kit (CSV+M), Power BI project (PBIP ZIP) into a Market Access-specific folder.
- Open HTML/CSV, compare values/units/scope/provenance to the preview, inspect ZIP/kit, and perform G2-G3 on this PBIP.

EVIDENCE

Capture table, chart, Audit trail, and evidence filenames.

22 M5: Market Access - repeat live request and REUSE

EVIDENCE
Record rendered plan ID/hash, repeat M1 with LLM enabled, and capture last extraction: llm.

- Confirm P-NUB/DE; coverage_pct_syn and reimbursement_status; empty months; snapshot; same owner OWN-001.
- Create plan. Expected REUSE, same original plan ID/hash and rendered frozen result; no new pending approval.
- Re-download CSV and compare source rows with initial export. Capture notice/IDs/Audit trail.

LAUNCH L1-L5

23 L1: Launch - exact LLM request

- Builder pack reports; Source domain Launch; requester DEMO-CHALLENGE-REQUESTER.

```
Show Nubeqa launch milestones in Japan in January 2026, including milestone ID, milestone type, and status.
```

- Enter request, enable LLM opt-in, click Extract intent; capture last extraction: llm and warnings before edits.
- Japan is JP, not Germany DE. Launch is an event dataset, not a sales or snapshot dataset.

LAUNCH L1-L5

24 L2: Launch - confirm event scope

- Products P-NUB; Markets JP.
- Source fields milestone_id, milestone_type, status.
- Months 2026-01-01; grain event. Explicit owner OWN-001; rationale Team Launch live LLM test.

EVIDENCE
Record any extraction mismatch before correcting the authoritative form.

EXPECTED
Expected Create plan enabled. Capture completed form with Japan and January visible.

25 L3: Launch - pending and separate approval

- Click Create plan; record ID/hash/revision. Fresh DB expected pending; if REUSE, follow L5 and record fresh-only checks not exercised.
- Choose DEMO-CHALLENGE-STEWARD in the approval block; click Approve. Expected API error 403, pending unchanged.
- Choose DEMO-CHALLENGE-APPROVER in the SAME block; click Approve. Expected approved.
- Do not change sidebar requester. Capture pending/refusal/approved and Audit trail.

26 L4: Launch - inspect events and all downloads

- Click Render frozen report. Expected rendered event table; do not require a numeric chart for categorical milestone fields.
- Check source event LM-0001 for P-NUB/JP with milestone_date 2026-01-05T00:00:00; preserve source milestone_type and status.
- No invented event totals, monthly sales, or snapshot trends.
- Download HTML, CSV, ZIP bundle, Power Query kit (CSV+M), Power BI project (PBIP ZIP) into a Launch-specific folder.
- Open HTML/CSV and compare to preview; inspect ZIP/kit; perform G2-G3 for Launch PBIP. Capture table/Audit trail/files.

27 L5: Launch - repeat live request and REUSE

EVIDENCE
Record rendered Launch plan ID/hash; repeat L1 with live LLM and capture last extraction: llm.

- Confirm P-NUB/JP; milestone_id, milestone_type, status; 2026-01-01; event; SAME owner OWN-001.
- Click Create plan. Expected REUSE of the same plan ID/hash and frozen result, no new pending approval.
- Compare re-downloaded CSV source rows; capture REUSE notice, IDs, and Audit trail.
- Do not silently replace January or Japan if extraction differs; record the mismatch.

28 G1: Inspect all five exports for EACH dataset

- HTML: open locally and check scope, native units, source rows, provenance, ownership, and audit information.
- CSV: check selected columns and source rows against the frozen preview. Do not let Excel's date/number formatting masquerade as changed source values.
- ZIP bundle: extract and inspect packaged report/evidence files. Power Query kit: inspect CSV, PowerQuery_Report.m, theme, and included instructions; this is a manual-import kit, NOT the PBIP project.
- PBIP ZIP is a different export: use G2-G3. Keep each domain's files in its own folder; automatic browser suffixes are not reliable domain labels.

EVIDENCE

Record five downloaded filenames per dataset. A visible download button is not proof that a usable file was downloaded.

29 G2: Power BI Desktop handoff - repeat four times

- Optional dependency: Power BI Desktop with PBIP/project support. If unavailable, record Desktop checks as not tested, not passed.
- For each Commercial, Supply, Market Access, and Launch PBIP ZIP, extract the FULL archive into its own folder.
- Keep GovernedReport.Report and GovernedReport.SemanticModel beside GovernedReport.pbip. Do not move only the .pbip file.
- Open GovernedReport.pbip in Desktop. Run Refresh on first load even if not prompted: this loads the embedded frozen snapshot, not new business data.

EVIDENCE

Record Desktop version and any warnings. Templates were previously checked with 2.159.1231.0; verify your own target version.

30 G3: Validate every Power BI page per dataset

- Open Dashboard, Source Evidence, and Ownership & Evidence; capture all three pages for EACH of the four datasets after Refresh.
- Compare scope, field names, values, units, owner, and source identifiers with that domain's frozen HTML/CSV and UI preview.
- Commercial: monthly sales/forecast source values. Supply/Market Access: snapshot values and categorical statuses, no invented dated trends. Launch: event identity/type/status, not sales totals.
- Measures may return BLANK for ambiguous multi-row contexts; do not demand summed/averaged source values. Investigate an unexpected blank in an unambiguous context.
- No PBIX export, service publish, live connection, or scheduled refresh is implemented. Power Query kit requires manual import and is not this project-opening test.

31 G4: Saved plans, reload, and URL resume

- Repeat for all four rendered plans: Workspace Saved plans, click Refresh, select the recorded plan ID.

EXPECTED
Expected: same plan ID/hash, rendered snapshot, Audit trail, and five download choices. Re-download CSV and compare source rows.

- Copy the UI URL containing ?plan_id=...; open in a new browser tab with the same requester identity.

EXPECTED
Expected report plan resumes in Builder pack reports with its domain and frozen result. A lifecycle plan resumes in Saved plans.

- Browser reload/new tab must not demand a fresh approval for an already-rendered plan. For persistence after restart, stop only your launcher and restart with the SAME recorded RUNTIME_DB_PATH and ports.
- Do NOT regenerate the GUID command when testing restart persistence; that would select a different empty DB.

32 G5: REUSE boundaries and scope changes

- REUSE uses the normalized report specification plus source fingerprint, including owner. Same wording alone is not enough; wording/rationale are not the deduplication key.
- On each rendered domain, reorder selected fields (and Commercial months), keeping the same selection set and owner; create again. Expected same original plan via REUSE.
- Changed-scope test: from Commercial C2, keep Germany/Nubeqa/OWN-001 but select ONLY net_sales_m_eur_syn (remove forecast); create.
- In this clean run that new scope must be pending, not reuse the original two-field plan. If previously approved, it may reuse its OWN different plan; record both IDs.
- A pending plan is not eligible for approved-plan REUSE. Do not claim metadata-only source-estate candidates are equivalent executable reports.

33 G6: Duplicate pending plans coalesce at approval

- Optional advanced UI test, before approving the new single-field scope from G5: create the SAME single-field scope again with the same owner, without approving either copy.

> **EXPECTED**
> Expected two different pending plan IDs, A and B. Record both; use Saved plans to select them.

- Approve A with DEMO-CHALLENGE-APPROVER. Then load B freshly from Saved plans and approve B.

> **EXPECTED**
> Expected second approval returns REUSE of A; B becomes superseded. Audit trail records the duplicate resolution; only the winning approved plan should be rendered.

- If the scope was already approved before this test, choose a new valid scope or a new isolated test DB; do not delete prior records.
- This tests duplicate handling, not a live multi-thread race. Stale-revision concurrency checks are in the automated suite.

34 G7: Rejection does not produce a report

- Use a fresh valid Commercial scope not yet approved: Nubeqa/Germany, ONLY forecast_m_eur_syn, January 2026, owner OWN-001.
- Create a pending plan; select DEMO-CHALLENGE-APPROVER in its approval block, then click Reject.

EXPECTED
Expected rejected status, rejection in Audit trail, and no Render frozen report button or export downloads for that rejected plan.

- Saved plans must preserve the rejected state. Creating the same rejected scope again creates a new pending plan unless an equivalent approved plan exists.
- Do not reject an already-approved baseline. Terminal approved/rejected decisions cannot simply be flipped; API checks are covered by G15.

35 G8: Finland scope refusal and identity reset

- After recording/saving current plan IDs, switch sidebar identity to DEMO-CHALLENGE-FI. Expected current plan/form identity state clears; stored plans are not deleted.
- In Commercial use the DE/Germany scope from C2 with active owner OWN-001 and click Create plan.

EXPECTED
Expected API error 403: requested markets outside your scope. Extracting text is not itself scope authorization; plan creation enforces it.

- Workspace Estate decisions as FI must also be refused (403), because estate records span all markets.
- Saved plans must not expose Germany/Japan-only plans to FI. Opening one of those saved URLs as FI must not reveal its frozen report.
- Restore DEMO-CHALLENGE-REQUESTER; use Saved plans to resume your original work. Simulated policy checks are not real login security.

36 G9: Required fields and authoritative confirmation

- On a new builder form, confirm Create plan is disabled with missing product, market, source fields, or explicit owner.
- For Commercial/Launch, missing required months also disables creation. Supply/Market Access have no period selector.
- Select a complete scope; owner must still be chosen explicitly. Capture disabled and enabled states.
- Intake warnings are visible, but they do not automatically prevent a user from explicitly choosing a different valid scope in the form.
- Do not report a manually corrected form as an accurate unedited LLM extraction. Preserve before/after evidence.

37 G10: Unsupported or incomplete intake

- In Commercial with LLM enabled, try: Show Nubeqa in Germany. Expect missing required scope to be surfaced; record any guessed fields as an extraction issue. Do not approve silently guessed scope.
- In Supply try: Show monthly Nubeqa supply risk in Germany for January through March 2026. Inspect unsupported-period/grain warnings; snapshot source must not become monthly data.
- In Commercial try: Show Nubeqa profit in Germany in January 2026. Profit is not a supported source field; record if extraction silently substitutes sales.
- In Launch try: Show Nubeqa milestones in Japan in December 2025. Out-of-coverage periods must not silently become January 2026.
- LLM wording varies; capture actual mode, warnings/questions, and proposed/confirmed form. Deterministic API invalid-scope checks are separately covered in G15.

38 G11: Offline path versus provider-outage fallback

- Offline control: uncheck LLM, run the C1 text through Extract intent; expected last extraction: offline. The separate Prefill commercial example fills a known form but is not an inference test.
- Live-provider outage: if an outage occurs, capture last extraction: offline_fallback_after_llm_unavailable and the visible offline-fallback warning; record live tests as blocked.
- Do not intentionally invalidate the shared token or disrupt another tester's network/process. For a deterministic simulated outage use the mocked fallback test in G15.
- Continuing with an explicitly confirmed offline form can test report governance, but does not complete any of the four live-LLM extraction checks.
- Restore the live checkbox for the next LLM test; checking it without a NEW extraction does not change the mode of a previous result.

39 G12: Estate migration proposal, never execution

- As DEMO-CHALLENGE-REQUESTER open Estate decisions. Filter assets by AST-0002; inspect Asset, Model/owner, KPIs, usage, lineage/downstream, quality/migration signals, and governance rules.
- Choose action migrate, active owner OWN-001, and a nonempty Rationale such as Team estate migration proposal test.
- Click Create lifecycle proposal. Expected pending; in its decision block test steward refusal, then choose DEMO-CHALLENGE-APPROVER and Approve.

EXPECTED
Expected Approved proposal; not executed. There is no report render or execution for a lifecycle plan; named release sign-offs remain proposed requirements.

- Reopen in Saved plans and inspect specification/Audit trail. The source estate is unchanged. Capture proposal evidence and non-execution message.

40 G13: Other estate actions and refusal paths

- Supported action choices: reuse, repair, consolidate, retire, migrate. These are ALL proposals only; estate action reuse is not report-plan REUSE.
- For repair/reuse, choose an asset, active owner, and nonempty rationale; inspect the resulting proposal and approval. Do not claim a source report was executed.
- For consolidate, inspect Consolidation candidates; explicitly choose a candidate target. Shared product/market/KPI metadata is NOT verified executable equivalence.
- When no consolidation candidates exist, the UI warns; do not invent a target. Invalid targets must be refused by the API.
- Migrate is restricted to supported source platform; retirement may be refused for downstream dependencies. Use G15 for exact deterministic rejection tests instead of guessing safe asset IDs.

EVIDENCE

Record the selected asset/action and actual outcome; no retire/migrate action modifies the workbook.

41 G14: Source requests, ownership, provenance

- Builder page: expand Source business requests (original rows), choose a Sample request, click Load request text.

EXPECTED
Expected original text loads into Describe the report. Context labels may contradict text and are not auto-truth; inspect them rather than silently trusting them.

- This is a separate exploratory intake test; do not replace the four exact baseline requests when collecting C1/S1/M1/L1 evidence.
- Inspect Source estate records in this product/market scope (metadata only), where available, and the plan Specification.
- Workspace Ownership & deployment explains simulated roles, proposed responsibility, and production gaps. Accountable owner is not enforced certification.
- Preserve native units, value_type, source-row references, and fingerprint. Exported local files have no access control after download.

42 G15: Automated API-only guardrail appendix

- These offline tests cover cases not fully exposed as clickable UI controls. They do NOT prove live provider availability and must be recorded separately from manual/live results.

```
$env:PYTHONPATH = 'src'
.\.venv\Scripts\python.exe -m unittest discover -s tests -p "test_challenge.py"
.\.venv\Scripts\python.exe -m unittest discover -s tests -p "test_pbip_dashboard.py"
.\.venv\Scripts\python.exe -m unittest discover -s tests -p "test_launcher_config.py"
```

- Challenge suite covers unknown IDs, wrong grain/months/owner, cross-domain fields, scope refusal, pending/rejected render refusal, stale hash/revision, source drift, duplicate approval coalescing, frozen resume, export packaging, and mocked LLM fallback.
- Never edit workbook/evidence to demonstrate source drift manually; tests isolate such scenarios. Legacy live_aitta scripts target different domains and are not this four-dataset acceptance flow.

43 G16: Troubleshooting without losing test evidence

- Busy ports: reuse the correct instance or choose a new API/UI pair. Changing ports alone does NOT create a fresh DB.
- Module errors: use the SAME venv Python for pip and launcher. Wrong UI: launch --mode challenge, not public/legacy.
- API unreachable: inspect launching terminal and artifacts/logs/api-8021.log / ui-8521.log (adjust port suffixes).
- LLM unavailable: verify approved credentials/provider/model privately, capture warning and mark live tests blocked; never publish token contents.
- Unexpected REUSE: compare original scope, owner, hash, and source fingerprint; fresh pending approval needs an unused scope/new isolated DB, not database deletion.
- Lost current UI plan after identity change: restore requester and reopen Saved plans; do not recreate blindly.

44 G17: Team completion record and shutdown

- For EACH dataset record: live extraction caption; unedited scope accuracy and corrections; confirmed owner/scope; pending or existing-REUSE branch; steward refusal; approval; rendered result; five downloaded files; repeat-request REUSE; saved-plan resume; optional Desktop results.
- Use result labels NOT RUN, PASS, FAIL, BLOCKED, or NOT APPLICABLE with an explanation. Never prefill success. Existing-state REUSE does not pass a fresh-approval test that was skipped.

EVIDENCE

Record shared G1-G16 results and evidence links separately. If Desktop/provider is unavailable, say so; do not relabel it a complete live end-to-end pass.

- Existing tests provide the cited synthetic spot-check values; your team must still compare current frozen output to source provenance. These instructions are not today's execution report.
- Stop your foreground demo with Ctrl+C in its launching terminal. Keep the recorded DB and downloads for handoff; do not delete someone else's processes or files.
- To return a dedicated test PowerShell to default DB selection later, remove only its process variable: Remove-Item Env:RUNTIME_DB_PATH. This does not delete a DB file; .env configuration still applies.